Findings of the 2009 Workshop on Statistical Machine Translation

Summary of the paper by Callison-Burch, et al. (2009)

Nakul Bhat
230962212

Ashwin Nair
230962354

Davasam Karthikeya
230962326

Adya Gangwal
230962084

Vasishta Nandipati
230962095

Meeti Mehta
230962052

School of Computer Engineering
Manipal Institute of Technology

April 22, 2026

2026-04-19

WSMT 2009

Findings of the 2009 Workshop on Statistical Machine Translation
Summary of the paper by Callison-Burch, et al. (2009)
Nakul Bhat 230962212
Ashwin Nair 230962354
Davasam Karthikeya 230962326
Adya Gangwal 230962084
Vasishta Nandipati 230962095
Meeti Mehta 230962052
School of Computer Engineering Manipal Institute of Technology
April 22, 2026

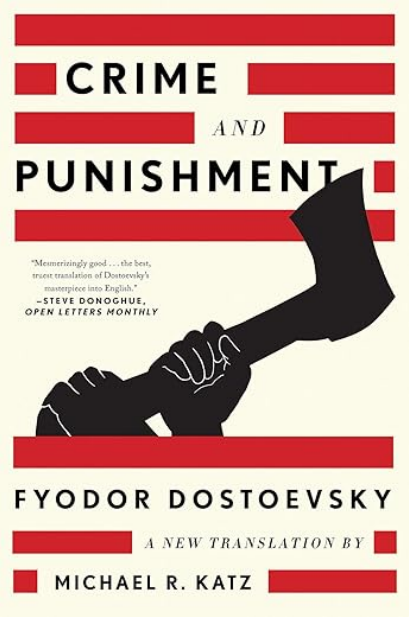

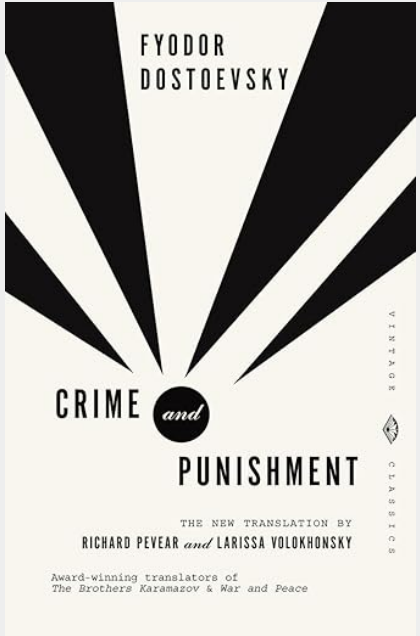

2026-04-19

WSMT 2009

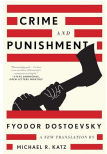

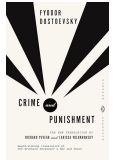

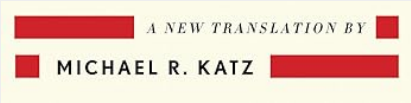
A NEW TRANSLATION BY
MICHAEL R. KATZ

THE NEW TRANSLATION BY
RICHARD PEVEAR and LARISSA VOLOKHONSKY
SLICE

A NEW TRANSLATION BY
MICHAEL R. KATZ

THE NEW TRANSLATION BY
RICHARD PEVEAR and LARISSA VOLOKHONSKY

WMT 2009 Overview

- Fourth Workshop on Statistical Machine Translation
- Shared task evaluation of machine translation systems

Goals of the Workshop

- Evaluate state-of-the-art machine translation
- Provide shared datasets and benchmarks
- Compare systems under common conditions
- Improve evaluation methodologies

WMT 2009 Overview
- Fourth Workshop on Statistical Machine Translation
- Shared task evaluation of machine translation systems
Goals of the Workshop
- Evaluate state-of-the-art machine translation
- Provide shared datasets and benchmarks
- Compare systems under common conditions
- Improve evaluation methodologies

WMT 2009 Overview

Translation task | System combination task | Evaluation task

Goals of the Workshop

- Evaluate state-of-the-art machine translation
- Provide shared datasets and benchmarks
- Compare systems under common conditions
- Improve evaluation methodologies

WMT 2009 Overview
Translation task System combination task Evaluation task
Goals of the Workshop
- Evaluate state-of-the-art machine translation
- Provide shared datasets and benchmarks
- Compare systems under common conditions
- Improve evaluation methodologies

Translation Task

- News domain translation
- English and five European languages

Training Data	Testing Data
Automated web crawler data, and existing datasets	Manually translated news articles.

Datasets Used

Translation Datasets

- Europarl Translation Corpus
- News Commentary Translation Corpus
- 10^9 Word Parallel Corpus

Language Datasets

- Europarl Language Model Corpus
- News Language Model Corpus

Translation Task
News domain translation
English and five European languages
Training Data
Testing Data
Automated web crawler data, and existing datasets
Manually translated news articles.
Datasets Used
Translation Datasets
Europarl Translation Corpus
News Commentary Translation Corpus
10^9 Word Parallel Corpus
Language Datasets
Europarl Language Model Corpus
News Language Model Corpus

System Combination Task

- Participants were allowed to use **Moses**, an open-source translation toolkit.
- Moses performed at par with the best submissions of the previous year.
- 22 groups from 20 institutions submitted; 12 groups used the Moses toolkit.
- 87 systems were submitted, and an additional 42 were submitted as secondary systems. 6 combination systems were also submitted.

System Combination

- Naïve Combination
- Multi-source combination

System Combination Task
Participants were allowed to use Moses, an open-source translation toolkit.
Moses performed at par with the best submissions of the previous year.
22 groups from 20 institutions submitted; 12 groups used the Moses toolkit.
87 systems were submitted, and an additional 42 were submitted as secondary systems. 6 combination systems were also submitted.
System Combination
Naïve Combination
Multi-source combination

Evaluation Task

- Outputs from select systems were evaluated by human annotators. These were split into two types:
  1. Ranked lists
  2. Edited translations
- Automated metrics also scored the translations.
- Both scores were compared to find the best automated metrics.

Evaluation Task
Outputs from select systems were evaluated by human annotators. These were split into two types:
Ranked lists
Edited translations
Automated metrics also scored the translations.
Both scores were compared to find the best automated metrics.

Next Presenter: Ashwin Nair

Human Evaluation

Ashwin Nair

Human Evaluation

Ashwin Nair

Statistics on Human Effort

- Manual evaluation done by:-
  - shared task participant
  - interested volunteers
  - small number of paid of annotators

Participation Statistics

- More than 160 people participated, with 100 people putting more than an hours worth of work, and 30 putting more than 4 hours.
- collectively 479 hours was spent.

Statistics on Human Effort
Manual evaluation done by:-
shared task participant
interested volunteers
small number of paid of annotators
Participation Statistics
More than 160 people participated, with 100 people putting more than an hours worth of work, and 30 putting more than 4 hours.
collectively 479 hours was spent.

Overview of Evaluation Strategy

Core Philosophy

Automatic metrics are imperfect, greater emphasis on human judgment

Three Primary Tasks

- Sentence Ranking: The official determinant of system quality.
- Editing MT Output: Assessing fluency and intelligibility.
- Acceptability Judging: Verifying the correctness of edited output.

Overview of Evaluation Strategy
Core Philosophy
Automatic metrics are imperfect, greater emphasis on human judgment
Three Primary Tasks
Sentence Ranking: The official determinant of system quality.
Editing MT Output: Assessing fluency and intelligibility.
Acceptability Judging: Verifying the correctness of edited output.

Task 1 — Sentence Ranking

The Official Metric

- The Process: Annotators are shown a source sentence and up to 5 system translations.
- The Instruction: ”Rank translations from Best to Worst (ties allowed).”
- Scoring: Systems are ranked based on how frequently they are judged ”better than or equal to” others.

Scale

Over 35,000 individual rankings were collected across all language pairs

Task 1 — Sentence Ranking
The Official Metric
- The Process: Annotators are shown a source sentence and up to 5 system translations.
- The Instruction: ”Rank translations from Best to Worst (ties allowed).”
- Scoring: Systems are ranked based on how frequently they are judged ”better than or equal to” others.
Scale
Over 35,000 individual rankings were collected across all language pairs

Task 2 — Editing MT Output

Measuring Fluency and Intelligibility

- Experimental Design: Judges see MT output without the source or reference translation.
- The Task: Edit the sentence to be as fluent as possible.

Three Outcomes

- Edited: Sentence was understandable but needed polish.
- No corrections needed: Sentence was already perfect.
- Unable to correct: MT output was too garbled to understand.

Task 2 — Editing MT Output
Measuring Fluency and Intelligibility
- Experimental Design: Judges see MT output without the source or reference translation.
- The Task: Edit the sentence to be as fluent as possible.
Three Outcomes
- Edited: Sentence was understandable but needed polish.
- No corrections needed: Sentence was already perfect.
- Unable to correct: MT output was too garbled to understand.

Task 2 — Editing MT Output

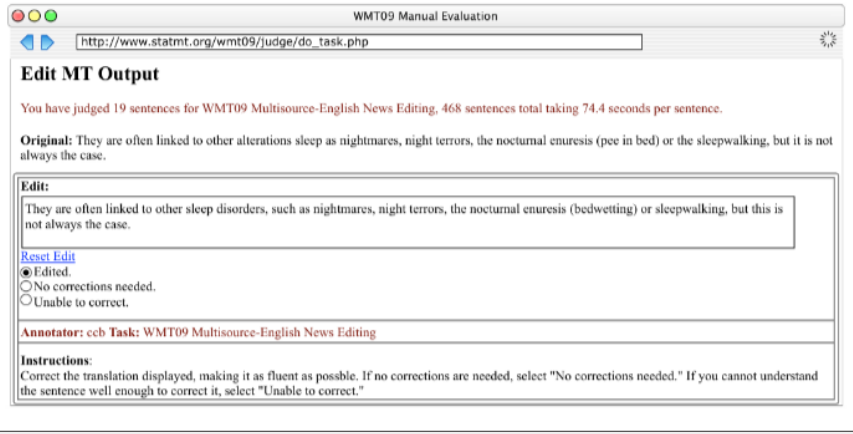

Figure: Interface for Editing MT Output

Task 2 — Editing MT Output

Figure: Interface for Editing MT Output

Task 3 — Judging Acceptability

Validation of the Edits

- The Setup: A second group of judges reviews the edited sentences alongside the original reference.
- The Question: Is this edited version a ”fully fluent and meaning-equivalent” alternative to the reference?

Outcome

A simple Yes/No indicator for over 31,000 edited items.

Task 3 — Judging Acceptability
Validation of the Edits
- The Setup: A second group of judges reviews the edited sentences alongside the original reference.
- The Question: Is this edited version a ”fully fluent and meaning-equivalent” alternative to the reference?
Outcome
A simple Yes/No indicator for over 31,000 edited items.

Task 3 — Judging Acceptability

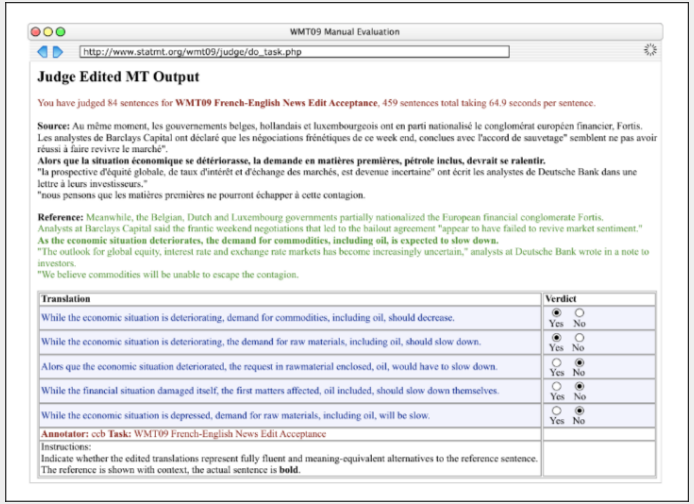

WMT09 Manual Evaluation

http://www.statmt.org/wmt09/judge/do_task.php

Judge Edited MT Output

You have judged 84 sentences for **WMT09 French-English News Edit Acceptance**, 459 sentences total taking 64.9 seconds per sentence.

Source: Au même moment, les gouvernements belges, hollandais et luxembourgeois ont **en parti** nationalisé le conglomérat européen financier, Fortis. Les analystes de Barclays Capital ont déclaré que les négociations frénétiques de ce week end, conclues avec l'accord de sauvetage" semblent ne pas avoir réussi à faire revivre le marché".
Alors que la situation économique se détériorasse, la demande en matières premières, pétrole inclus, devrait se ralentir.
"la prospective d'équité globale, de taux d'intérêt et d'échange des marchés, est devenue incertaine" ont écrit les analystes de Deutsche Bank dans une lettre à leurs investisseurs."
"nous pensons que les matières premières ne pourront échapper à cette contagion.

Reference: Meanwhile, the Belgian, Dutch and Luxembourg governments partially nationalized the European financial conglomerate Fortis.
Analysts at Barclays Capital said the frantic weekend negotiations that led to the bailout agreement "appear to have failed to revive market sentiment."
As the economic situation deteriorates, the demand for commodities, including oil, is expected to slow down.
"The outlook for global equity, interest rate and exchange rate markets has become increasingly uncertain," analysts at Deutsche Bank wrote in a note to investors.
"We believe commodities will be unable to escape the contagion.

Translation	Verdict
While the economic situation is deteriorating, demand for commodities, including oil, should decrease.	◉ Yes ○ No
While the economic situation is deteriorating, the demand for raw materials, including oil, should slow down.	◉ Yes ○ No
Alors que the economic situation deteriorated, the request in rawmaterial enclosed, oil, would have to slow down.	○ Yes ◉ No
While the financial situation damaged itself, the first matters affected, oil included, should slow down themselves.	○ Yes ◉ No
While the economic situation is depressed, demand for raw materials, including oil, will be slow.	○ Yes ◉ No
Annotator: ccb **Task:** WMT09 French-English News Edit Acceptance	
Instructions: Indicate whether the edited translations represent fully fluent and meaning-equivalent alternatives to the reference sentence. The reference is shown with context, the actual sentence is **bold**.	

Figure: Interface for Editing MT Output

Task 3 — Judging Acceptability

Figure: Interface for Editing MT Output

Quality Control – Annotator Agreement

Measuring Reliability (Kappa Coefficient K)

$$K = (P(A) - P(E))(1 - P(E))$$

- P(A) - the proportion of times that the annotators agree
- P(E) - the proportion of time that they would agree by chance.

- Intra-annotator agreement: How often a judge agrees with themselves (Substantial: K = 0.56 – 0.73).
- Inter-annotator agreement: How often different judges agree (Fair/Moderate: K = 0.32 – 0.54).

Optimization

The study found that removing the 33 worst annotators (those with the lowest agreement) improved overall data reliability while retaining 60% of the judgments

Quality Control – Annotator Agreement
Measuring Reliability (Kappa Coefficient K)
$K = (P(A) - P(E))(1 - P(E))$
P(A) - the proportion of times that the annotators agree
P(E) - the proportion of time that they would agree by chance.
Intra-annotator agreement: How often a judge agrees with themselves (Substantial: K = 0.56 – 0.73).
Inter-annotator agreement: How often different judges agree (Fair/Moderate: K = 0.32 – 0.54).
Optimization
The study found that removing the 33 worst annotators (those with the lowest agreement) improved overall data reliability while retaining 60% of the judgments

Task 3 — Judging Acceptability

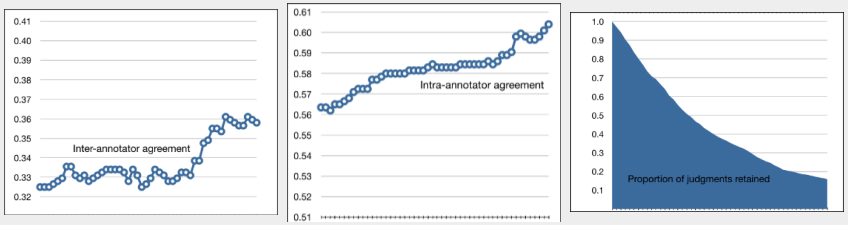

Figure: Dropping the first 50 scores from each evaluator.

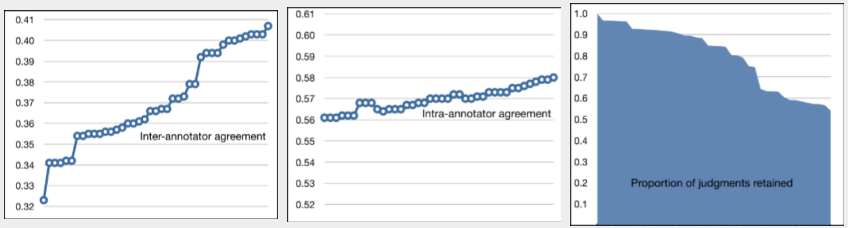

Figure: Dropping the worst 40 evaluators.

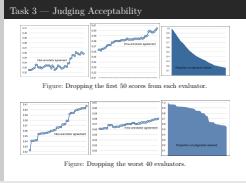

Key Takeaways

- Human-Centric: WMT09 prioritizes human ranking over BLEU scores to validate systems.
- Scale: Massive crowdsourced effort from participants, volunteers, and paid staff.
- Innovation: Introduced "editing-based" evaluation to move beyond simple ranking.
- Transparency: All human judgments are made publicly available for future research.

Key Takeaways
- Human-Centric: WMT09 prioritizes human ranking over BLEU scores to validate systems.
- Scale: Massive crowdsourced effort from participants, volunteers, and paid staff.
- Innovation: Introduced "editing-based" evaluation to move beyond simple ranking.
- Transparency: All human judgments are made publicly available for future research.

Next Presenter: Davasam Karthikeya

Translation Results

Davasam Karthikeya

Translation Results

Davasam Karthikeya

Translation Task Results

In this section, we will try to answer these three questions:

- Which systems produced the best translation quality for each language pair?
- Did the system combinations produce better translations than individual systems?
- Which of the systems that used only the provided training materials produced the best translation quality?

Translation Task Results
In this section, we will try to answer these three questions:
- Which systems produced the best translation quality for each language pair?
- Did the system combinations produce better translations than individual systems?
- Which of the systems that used only the provided training materials produced the best translation quality?

Which MT systems are best?

How "Best" is Defined?

- Based on human evaluation (sentence ranking)
- Systems ranked by:
- How often their translations are better than or equal to others
- They used Sign Test (p 0.1)

Pair	Performance Range	Top systems (C)	Top systems (NC)
French → English	0.76 - 0.34	DCU, JHU	Google, LIMSI
English → French	0.73 - 0.38	UKA, DCU	LIUM-SYSTRAN, Google

Which MT systems are best?
How "Best" is Defined?
- Based on human evaluation (sentence ranking)
- Systems ranked by:
- How often their translations are better than or equal to others
- They used Sign Test (p 0.1)

Pair	Performance Range	Top systems (C)	Top systems (NC)
French → English	0.76 - 0.34	DCU, JHU	Google, LIMSI
English → French	0.73 - 0.38	UKA, DCU	LIUM-SYSTRAN, Google

Which MT systems are best?

Pair	Performance Range	Top systems (C)	Top systems (NC)
German → English	0.66 - 0.33	STUTTGART, UEDIN	RBMT5, US-AAR
English → German	0.66 - 0.43	UKA, UEDIN	RBMT2, RBMT3
Czech → English	0.75 - 0.51	UEDIN, CU-BOJAR	Google
Hungarian → English	0.75 - 0.45	UMD, UEDIN	MORPHO

Which MT systems are best?

Pair	Performance Range	Top systems (C)	Top systems (NC)
German → English	0.66 - 0.33	STUTTGART, UEDIN	RBMT5, US-AAR
English → German	0.66 - 0.43	UKA, UEDIN	RBMT2, RBMT3
Czech → English	0.75 - 0.51	UEDIN, CU-BOJAR	Google
Hungarian → English	0.75 - 0.45	UMD, UEDIN	MORPHO

Which MT systems are best?

Pair	Performance Range	Top systems (C)	Top systems (NC)
Spanish → English	0.70 - 0.37	TALP-UPC, WEDIN	Google, RBMT1
English → Spanish	0.66 - 0.48	UEDIN, NUS	RBMT3, Google
English → Czech	0.67 - 0.48	CU-BOJAR, WEDIN	PCTRANS, EURO-TRANXP

Which MT systems are best?

Pair	Performance Range	Top systems (C)	Top systems (NC)
Spanish → English	0.70 - 0.37	TALP-UPC, WEDIN	Google, RBMT1
English → Spanish	0.66 - 0.48	UEDIN, NUS	RBMT3, Google
English → Czech	0.67 - 0.48	CU-BOJAR, WEDIN	PCTRANS, EURO-TRANXP

Are system combinations better than single systems?

Understandability

- Measures how often translated output is clear and meaningful after editing
- Based on human judgment (Yes/No) on edited translations
- Reflects real-world usability of translations

Results:

- 50% acceptable translations (French - English, German - English, English - Spanish)
- 40% acceptable translations (Spanish - English)
- 30% acceptable translations (English - Czech)
- 20% acceptable translations (Czech - English, Hungarian - English)
- Even top systems are only correct half the time
- Large variation across language pairs

Are system combinations better than single systems?
Understandability
- Measures how often translated output is clear and meaningful after editing
- Based on human judgment (Yes/No) on edited translations
- Reflects real-world usability of translations
Results:
- 50% acceptable translations (French - English, German - English, English - Spanish)
- 40% acceptable translations (Spanish - English)
- 30% acceptable translations (English - Czech)
- 20% acceptable translations (Czech - English, Hungarian - English)
- Even top systems are only correct half the time
- Large variation across language pairs

Takeaways

- Human evaluation is the gold standard
- Commercial systems (Google, RBMT) perform strongly
- Performance drops for complex/hard language pairs
- Many systems perform similarly → hard to declare a clear winner
- System combinations doesn't guarantee improvement
- Machine Translation quality is still far from perfect

Takeaways
- Human evaluation is the gold standard
- Commercial systems (Google, RBMT) perform strongly
- Performance drops for complex/hard language pairs
- Many systems perform similarly → hard to declare a clear winner
- System combinations doesn't guarantee improvement
- Machine Translation quality is still far from perfect

Next Presenter: Adya Gangwal

Shared Evaluation

Adya Gangwal

Shared Evaluation

Adya Gangwal

Shared Evaluation Task Overview

- Goal: Evaluate automatic MT evaluation metrics
- Compare metrics against human judgments
- Builds on previous Metrics-MATR evaluation (NIST)

Table: Automatic Metrics Evaluated by Type

Metric Type	Specific Metrics Evaluated
N-gram based	BLEU, METEOR
Edit-distance	TER, TERp
Similarity-based	MaxSim
Semantic	RTE, SemPOS
Combined	ULC
POS-based	wpF, wpBLEU

Shared Evaluation Task Overview
- Goal: Evaluate automatic MT evaluation metrics
- Compare metrics against human judgments
- Builds on previous Metrics-MATR evaluation (NIST)
Table: Automatic Metrics Evaluated by Type
Metric Type | Specific Metrics Evaluated
N-gram based | BLEU, METEOR
Edit-distance | TER, TERp
Similarity-based | MaxSim
Semantic | RTE, SemPOS
Combined | ULC
POS-based | wpF, wpBLEU

System-Level Correlation

- Compare metric rankings with human rankings
- Use Spearman rank correlation (ρ)

$$\rho = 1 - \frac{6 \sum d_i^2}{n(n^2 - 1)}$$

where,

d_i	Difference between the rank for system i
n	Number of systems

- Convert system scores into ranks
- Measure similarity with human judgments

System-Level Correlation
- Compare metric rankings with human rankings
- Use Spearman rank correlation (ρ)
$$\rho = 1 - \frac{6 \sum d_i^2}{n(n^2 - 1)}$$
where,
d_i Difference between the rank for system i
n Number of systems
- Convert system scores into ranks
- Measure similarity with human judgments

Sentence-Level Consistency

- Compare metrics at sentence level
- Use pairwise system comparisons
- Check agreement with human judgments
- Report percentage of consistent decisions

Sentence-Level Consistency
- Compare metrics at sentence level
- Use pairwise system comparisons
- Check agreement with human judgments
- Report percentage of consistent decisions

Next Presenter: Vasishta Nandipati

Evaluation Results

Vasishta Nandipati

Evaluation Results

Vasishta Nandipati

System-Level Correlation

Method

Spearman's rank correlation (ρ).
$\rho = 1$ means perfect match with humans.
$\rho =$ -1 means perfect reverse match with humans.

Results

- For non-English targets: metrics relying on English tools (WordNet, POS taggers) fail. TERp and wpBleu are best there.
- English-German: nearly all metrics had negative correlation.

Metric	Avg rho
ULC	0.83
MaxSim	0.80
RTE (absolute)	0.79
Meteor-rank	0.75
BLEU	0.47
TER / TERp	negative

System-Level Correlation
Method
Spearman's rank correlation (ρ).
$\rho = 1$ means perfect match with humans.
$\rho =$ -1 means perfect reverse match with humans.
Results
For non-English targets: metrics relying on English tools (WordNet, POS taggers) fail. TERp and wpBleu are best there.
English-German: nearly all metrics had negative correlation.
Metric Avg rho
ULC 0.83
MaxSim 0.80
RTE (absolute) 0.79
Meteor-rank 0.75
BLEU 0.47
TER / TERp negative

Sentence-Level Consistency

Method

for every pair of systems on one sentence, did the metric agree with the human? Random baseline = 50%.

Most metrics barely beat – or fall below – 50%.

Into English best

Metric	Score
ULC	0.54
RTE	0.53
wpF	0.52

Out of English best

Metric	Score
wpF	0.51
wcd6p4er	0.50

- TER / TERp: fell BELOW 50% – worse than random.

Sentence-Level Consistency
Method
for every pair of systems on one sentence, did the metric agree with the human? Random baseline = 50%.
Most metrics barely beat – or fall below – 50%.
Into English best
Metric Score
ULC 0.54
RTE 0.53
wpF 0.52
Out of English best
Metric Score
wpF 0.51
wcd6p4er 0.50
TER / TERp: fell BELOW 50% – worse than random.

Sentence-Level Consistency (Tables 10-11)

The "prior" trick

using a metric's system-level score as a fixed prior for every sentence it produces jumps consistency to 0.57-0.61, near Oracle (0.62).

Insight

which system produced the sentence predicts quality better than analyzing the sentence itself.

Sentence-Level Consistency (Tables 10-11)
The "prior" trick
using a metric's system-level score as a fixed prior for every sentence it produces jumps consistency to 0.57-0.61, near Oracle (0.62).
Insight
which system produced the sentence predicts quality better than analyzing the sentence itself.

Key Takeaways

- ULC best predicts human judgment at system level. BLEU is overused and underperforms.
- For non-English targets, TERp and wpBleu hold up; most others break down.
- Sentence-level evaluation is extremely hard – most metrics are near coin-flip.
- Using system-level performance as a prior significantly improves sentence-level consistency – a clear direction for future work.

Key Takeaways
- ULC best predicts human judgment at system level. BLEU is overused and underperforms.
- For non-English targets, TERp and wpBleu hold up; most others break down.
- Sentence-level evaluation is extremely hard – most metrics are near coin-flip.
- Using system-level performance as a prior significantly improves sentence-level consistency – a clear direction for future work.

Next Presenter: Meeti Mehta

Conclusion

Meeti Mehta

Summary and Task Simplification

- 22 research groups from 20 institutions
- 7 commercial systems and Google's online statistical MT system
- This time the focus was completely on news articles
- The unfair advantage of domain-specific training data? Eliminated.

Summary and Task Simplification
- 22 research groups from 20 institutions
- 7 commercial systems and Google's online statistical MT system
- This time the focus was completely on news articles
- The unfair advantage of domain-specific training data? Eliminated.

Human Evaluation Insights

- Introduction of an editing task to measure comprehensibility
- Learnt a lot about human reliability
  - Tried dropping first 50 learning judgements
  - Tried removing judges with the worst agreement

Human Evaluation Insights
- Introduction of an editing task to measure comprehensibility
- Learnt a lot about human reliability
  - Tried dropping first 50 learning judgements
  - Tried removing judges with the worst agreement

Automatic Metric Performance

- ULC metric correlated best for translations into English
- TERp and wpBleu performed best for the other languages
- Sentence level prediction proved to be challenging
- Using system-level scores as prior probabilities improved sentence-level scoring

Automatic Metric Performance
- ULC metric correlated best for translations into English
- TERp and wpBleu performed best for the other languages
- Sentence level prediction proved to be challenging
- Using system-level scores as prior probabilities improved sentence-level scoring